

Dimensionality Reduction: Linear projective techniques

Minería de Datos (M1966)

Steven Van Vaerenbergh
Depto. de Matemáticas, Estadística y Computación
Universidad de Cantabria

Curso 2024-25

Maître Universitario Oficial **Data Science**



Data format

Table with rows and columns:

columns
(features)

rows
(instances)

	Name	Number	Position	Age	Height	Weight	College	Salary
0	Avery Bradley	0.0	PG	25.0	6-2	180.0	Texas	7730337.0
1	Jae Crowder	99.0	SF	25.0	6-6	235.0	Marquette	6796117.0
3	R.J. Hunter	28.0	SG	22.0	6-5	185.0	Georgia State	1148640.0
6	Jordan Mickey	55.0	PF	21.0	6-8	235.0	LSU	1170960.0
7	Kelly Olynyk	41.0	C	25.0	7-0	238.0	Gonzaga	2165180.0
8	Terry Rozier	12.0	PG	22.0	6-2	190.0	Louisville	1824360.0
9	Marcus Smart	36.0	PG	22.0	6-4	220.0	Oklahoma State	3431040.0
10	Jared Sullinger	7.0	C	24.0	6-9	260.0	Ohio State	2569260.0
...
444	Alec Burks	10.0	SG	24.0	6-6	214.0	Colorado	9463484.0
446	Derrick Favors	15.0	PF	24.0	6-10	265.0	Georgia Tech	12000000.0
448	Gordon Hayward	20.0	SF	26.0	6-8	226.0	Butler	15409570.0
449	Rodney Hood	5.0	SG	23.0	6-8	206.0	Duke	1348440.0
451	Chris Johnson	23.0	SF	26.0	6-6	206.0	Dayton	981348.0
452	Trey Lyles	41.0	PF	20.0	6-10	234.0	Kentucky	2239800.0
453	Shelvin Mack	8.0	PG	26.0	6-3	203.0	Butler	2433333.0
456	Jeff Withey	24.0	C	26.0	7-0	231.0	Kansas	947276.0

Data format

Elimination of some features → **Feature Selection**.

columns
(features)

rows
(instances)

	Name	Number	Position	Age	Height	Weight	College	Salary
0	Avery Bradley	0.0	PG	25.0	6-2	180.0	Texas	7730337.0
1	Jae Crowder	99.0	SF	25.0	6-6	235.0	Marquette	6796117.0
3	R.J. Hunter	28.0	SG	22.0	6-5	85.0	Georgia State	1148640.0
6	Jordan Mickey	55.0	PF	21.0	6-8	235.0	LSU	1170960.0
7	Kelly Olynyk	41.0	C	25.0	7-0	238.0	Gonzaga	2165160.0
8	Terry Rozier	12.0	PG	22.0	6-2	180.0	Louisville	1824360.0
9	Marcus Smart	36.0	PG	21.0	6-4	220.0	Oklahoma State	3431040.0
10	Jared Sullinger	7.0	C	24.0	6-9	269.0	Ohio State	2569260.0
...
444	Alec Burks	10.0	SG	24.0	6-6	210.0	Colorado	9463484.0
446	Derrick Favors	15.0	PF	24.0	6-10	260.0	Georgia Tech	12000000.0
448	Gordon Hayward	20.0	SF	26.0	6-8	216.0	Butler	15409570.0
449	Rodney Hood	5.0	SG	23.0	6-8	206.0	Duke	1348440.0
451	Chris Johnson	23.0	SF	26.0	6-6	206.0	Dayton	981348.0
452	Trey Lyles	41.0	PF	20.0	6-10	234.0	Kentucky	2239800.0
453	Shelvin Mack	8.0	PG	26.0	6-3	203.0	Butler	2433333.0
456	Jeff Withey	24.0	C	26.0	7-0	231.0	Kansas	947276.0

Data format

Converting the original d features to a smaller set of r **new** features \rightarrow **Dimensionality Reduction**.

	a	b	c	d	e	f	g	h	i	j	k	l	m
0	0.0	0.0	12.0	13.0	5.0	0.0	0.0	0.0	0.0	0.0	11.0	16.0	9.0
1	0.0	0.0	4.0	15.0	12.0	0.0	0.0	0.0	0.0	3.0	16.0	15.0	14.0
2	0.0	7.0	15.0	13.0	1.0	0.0	0.0	0.0	8.0	13.0	6.0	15.0	4.0
3	0.0	0.0	1.0	11.0	0.0	0.0	0.0	0.0	0.0	0.0	7.0	8.0	0.0
4	0.0	12.0	10.0	0.0	0.0	0.0	0.0	0.0	0.0	14.0	16.0	18.0	14.0
5	0.0	0.0	12.0	13.0	0.0	0.0	0.0	0.0	0.0	5.0	16.0	8.0	0.0
6	0.0	7.0	8.0	13.0	16.0	15.0	1.0	0.0	0.0	7.0	7.0	4.0	11.0
7	0.0	9.0	14.0	8.0	1.0	0.0	0.0	0.0	0.0	12.0	14.0	14.0	12.0
8	0.0	11.0	12.0	0.0	0.0	0.0	0.0	0.0	2.0	16.0	16.0	16.0	13.0
9	0.0	1.0	9.0	15.0	11.0	0.0	0.0	0.0	0.0	11.0	16.0	8.0	14.0
10	0.0	0.0	0.0	14.0	13.0	1.0	0.0	0.0	0.0	0.0	5.0	16.0	16.0
11	0.0	5.0	12.0	1.0	0.0	0.0	0.0	0.0	0.0	15.0	14.0	7.0	0.0
12	2.0	9.0	15.0	14.0	9.0	3.0	0.0	0.0	4.0	13.0	8.0	9.0	16.0
13	0.0	0.0	8.0	15.0	1.0	0.0	0.0	0.0	0.0	1.0	14.0	13.0	1.0
14	5.0	12.0	13.0	16.0	16.0	2.0	0.0	0.0	11.0	16.0	15.0	8.0	4.0
15	0.0	0.0	8.0	15.0	1.0	0.0	0.0	0.0	0.0	0.0	12.0	14.0	0.0
16	0.0	1.0	8.0	15.0	10.0	0.0	0.0	0.0	3.0	13.0	15.0	14.0	14.0
17	0.0	10.0	7.0	13.0	9.0	0.0	0.0	0.0	0.0	9.0	10.0	12.0	15.0
18	0.0	6.0	14.0	4.0	0.0	0.0	0.0	0.0	0.0	11.0	16.0	10.0	0.0



θ	p	q	r
0	6.714282	4.263453	-3.698843
1	10.808431	-3.605318	-5.503636
2	-5.706837	1.612461	4.903340
3	7.967714	13.080329	0.804771
4	-13.494903	-2.883923	-9.851133
5	4.028233	11.200699	4.443256
6	7.745553	-10.512425	8.192260
7	-8.213388	-1.709213	-4.202378
8	-15.189400	-2.985254	-7.007185
9	3.892792	6.433339	0.217544
10	16.222026	-4.871032	-7.677208
11	-12.432569	7.402058	0.393797
12	-1.928984	1.383136	2.937210
13	5.864160	11.300092	0.921414
14	-3.750759	-10.217692	1.487572
15	6.721941	12.206578	0.701127
16	2.723118	-6.903958	-1.785841
17	1.891296	-7.808666	-3.138556
18	-10.273747	8.265386	2.453369

Why reduce dimensionality?

- The volume and the dimensionality of the data $\mathbf{x} \in \mathbb{R}^d$ in machine learning (ML) applications grows constantly.


Why reduce dimensionality?

- ▶ The volume and the dimensionality of the data $\mathbf{x} \in \mathbb{R}^d$ in machine learning (ML) applications grows constantly.
- ▶ ML techniques are not effective in high-dimensional spaces → **Curse of Dimensionality**




Why reduce dimensionality?

- ▶ The volume and the dimensionality of the data $\mathbf{x} \in \mathbb{R}^d$ in machine learning (ML) applications grows constantly.
- ▶ ML techniques are not effective in high-dimensional spaces → **Curse of Dimensionality**




Why reduce dimensionality?

- ▶ The volume and the dimensionality of the data $\mathbf{x} \in \mathbb{R}^d$ in machine learning (ML) applications grows constantly.
- ▶ ML techniques are not effective in high-dimensional spaces → **Curse of Dimensionality**



Why reduce dimensionality?


- ▶ The volume and the dimensionality of the data $\mathbf{x} \in \mathbb{R}^d$ in machine learning (ML) applications grows constantly.
- ▶ ML techniques are not effective in high-dimensional spaces → **Curse of Dimensionality**



Examples: Datos biomédicos, datasets de IoT (p.ej. series temporales de sensores), etc.

Why reduce dimensionality?


- The intrinsic dimension of the data of interest can be much less than the extrinsic data of the observation space or feature space.




The digit “6” can be represented by a small set of parameters.

Why reduce dimensionality?

- The intrinsic dimension of the data of interest can be much less than the extrinsic data of the observation space or feature space.



The digit “6” can be represented by a small set of parameters.



Other possible advantages of working in a lower dimensional data space:

- ▶ **Visualization**: Projection in 2D or 3D space.
- ▶ **Compression**: Reduction of the storage requirements while maintaining the possibility to recover the original data.
- ▶ **Noise reduction**: Projection onto a subspace or *manifold* in which the data of interest reside.
- ▶ **Convergence**: Dimensionality reduction improves the convergence of regression algorithms. (It reduces multicollinearity.)

Dimensionality reduction

Problem

Given n patterns or input vectors of dimension d , $\mathbf{x}_i \in \mathbb{R}^d$ ($i = 1, \dots, n$) and a desired output space dimension, $r < d$, the problem consists in finding a transformation

$$\mathbf{x}_i \in \mathbb{R}^d \longrightarrow \mathbf{y}_i \in \mathbb{R}^r,$$

that preserves/optimizes some characteristic of the data (e.g. variance, distances, inter-class separation).

Classification of techniques

- ▶ **Linear techniques** (Projective): A linear transformation between the input space of dimension d and the output dimension space of dimension $r < d$

$$\mathbf{y}_i = \mathbf{P}^T \mathbf{x}_i, \quad \mathbf{P} \in \mathbb{R}^{d \times r}$$

- ▶ Principal Component Analysis (PCA) and Linear Discriminant Analysis (LDA).

Classification of techniques

- ▶ **Linear techniques** (Projective): A linear transformation between the input space of dimension d and the output dimension space of dimension $r < d$

$$\mathbf{y}_i = \mathbf{P}^T \mathbf{x}_i, \quad \mathbf{P} \in \mathbb{R}^{d \times r}$$

- ▶ Principal Component Analysis (PCA) and Linear Discriminant Analysis (LDA).
- ▶ **Non-linear techniques**: Techniques that aim to model/extract the nonlinear *manifold* on which the data reside.
 - ▶ Multidimensional Scaling (MDS), Isomap, Locally Linear Embedding (LLE), Stochastic Neighbor Embedding (SNE).
 - ▶ Kernel-based methods: Kernel PCA, etc.


PCA: Introduction

Example: Direction of maximum variance

- ▶ Given a data set $\mathbf{x}_i \in \mathbb{R}^d$ ($i = 1, \dots, n$) with zero mean such that

$$\mathbf{x}_i = \theta_i \mathbf{u}, \quad \sum_{i=1}^n \theta_i = 0$$

where \mathbf{u} is a unit norm vector in \mathbb{R}^d . In other words, all data points \mathbf{x}_i lie on a line spanned by a unit vector \mathbf{u} .




- ▶ The projection of \mathbf{x}_i onto \mathbf{u} is $y_i = \mathbf{u}^T \mathbf{x}_i = \theta_i$.
- ▶ The variance of the entire projected data set is

$$\sigma_u^2 = \frac{1}{n} \sum_{i=1}^n y_i^2 = \frac{1}{n} \sum_{i=1}^n \theta_i^2$$

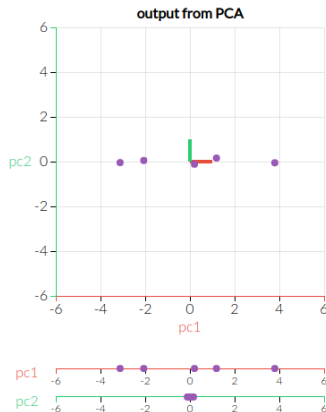

- ▶ If we project the \mathbf{x}_i onto any other unit norm vector $\mathbf{v} \neq \mathbf{u}$, the variance will be smaller

$$\sigma_v^2 = \frac{1}{n} \sum_{i=1}^n \theta_i^2 \langle \mathbf{v}, \mathbf{u} \rangle^2 < \sigma_u^2$$

Example: 2D data





2D example



Source and interactive example: <http://setosa.io/ev/principal-component-analysis/>

3D example



PCA theory (1/3)

- ▶ Given a data set $\mathbf{x}_i \in \mathbb{R}^d$ ($i = 1, \dots, n$) with zero mean (if the mean is not zero, we subtract the sample mean $\mathbf{m}_x = \frac{1}{n} \sum_i \mathbf{x}_i$)

$$\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_n]^T \in \mathbb{R}^{n \times d}$$

- ▶ The **sample covariance** matrix (dimensions $d \times d$) is

$$\hat{\mathbf{C}}_x = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i \mathbf{x}_i^T = \frac{1}{n} \mathbf{X}^T \mathbf{X}$$

- ▶ If we choose a direction $\mathbf{v} \in \mathbb{R}^d$ such that $\|\mathbf{v}\|^2 = 1$, the variance of the data projected onto this direction is

$$\sigma_v^2 = \mathbf{v}^T \hat{\mathbf{C}}_x \mathbf{v}$$

This measures how “spread out” the data is along the vector \mathbf{v} .

PCA theory (2/3)

- **Decompose $\hat{\mathbf{C}}_x$ into eigenvectors and eigenvalues:**

$$\hat{\mathbf{C}}_x = \mathbf{U}\mathbf{\Sigma}\mathbf{U}^T$$

where $\mathbf{\Sigma} = \text{diag}(\sigma_1^2, \dots, \sigma_d^2)$ with $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_d$

- Since $\mathbf{U} = [\mathbf{u}_1 \ \dots \ \mathbf{u}_d]$ forms a basis of \mathbb{R}^d , \mathbf{v} can be expanded as

$$\mathbf{v} = \sum_{i=1}^n \alpha_i \mathbf{u}_i, \quad \text{where} \quad \sum_{i=1}^n \alpha_i^2 = 1$$

and the variance of the projection is

$$\sigma_v^2 = \sum_i \alpha_i \sigma_i^2$$

PCA theory (3/3)

- ▶ The projection of maximum variance is found by solving

$$\underset{\alpha_i}{\text{maximize}} \sum_i \alpha_i \sigma_i^2 \quad \text{s.t.} \quad \sum_{i=1}^d \alpha_i^2 = 1$$

whose solution is $\alpha_1^* = 1, \alpha_2^* = \dots = \alpha_d^* = 0$

- ▶ The **direction that maximizes the variance is the principal eigenvector** of $\hat{\mathbf{C}}_x$

$$\mathbf{v} = \mathbf{u}_1$$


- ▶ The **first principal component** (1D projection) is

$$y_i = \mathbf{u}_1^T \mathbf{x}_i$$


- ▶ The proportion of **variance retained** by y_i is $\frac{\sigma_1^2}{\sum_{i=1}^d \sigma_i^2}$

Example: PCA on 2D data

Data (blue) and principal directions (red).



Projection onto the first principal direction.



► Python code in `example_1_pca_toy.ipynb`

Extension to r principal components

- ▶ Since the eigenvectors form an orthogonal basis and the eigenvalues are in descending order, the r PCA directions are

$$\mathbf{U}_r = [\mathbf{u}_1 \quad \dots \quad \mathbf{u}_r] \in \mathbb{R}^{d \times r}$$

- ▶ The data with reduced dimension are (principal components)

$$\mathbf{y}_i = \mathbf{U}_r^T \mathbf{x}_i \in \mathbb{R}^r \Rightarrow \mathbf{Y} = \mathbf{XU}_r$$

- ▶ The proportion of variance retained by y_i is now $\frac{\sum_{i=1}^r \sigma_i^2}{\sum_{i=1}^d \sigma_i^2}$.

Extension to r principal components

- ▶ Since the eigenvectors form an orthogonal basis and the eigenvalues are in descending order, the r PCA directions are

$$\mathbf{U}_r = [\mathbf{u}_1 \quad \dots \quad \mathbf{u}_r] \in \mathbb{R}^{d \times r}$$

- ▶ The data with reduced dimension are (principal components)

$$\mathbf{y}_i = \mathbf{U}_r^T \mathbf{x}_i \in \mathbb{R}^r \Rightarrow \mathbf{Y} = \mathbf{XU}_r$$

- ▶ The proportion of variance retained by \mathbf{y}_i is now $\frac{\sum_{i=1}^r \sigma_i^2}{\sum_{i=1}^d \sigma_i^2}$.
- ▶ We often choose r so that we retain, e.g., 90 % or 95 % of total variance.
- ▶ We can approximate the original data via $\hat{\mathbf{X}} = \mathbf{YU}_r^T$.

Python examples

- ▶ 3D example: `example_2_pca_3D.ipynb`
- ▶ Exercise 1:
`exercise_1_iris_visualization_PCA.ipynb`

PCA decorrelates the projected data

The r principal components of the vector $\mathbf{y}_i = \mathbf{U}_r^T \mathbf{x}_i$ are uncorrelated

$$\begin{aligned} E [\mathbf{y}_i \mathbf{y}_i^T] &= E [\mathbf{U}_r^T \mathbf{x}_i \mathbf{x}_i^T \mathbf{U}_r] = \mathbf{U}_r^T E [\mathbf{x}_i \mathbf{x}_i^T] \mathbf{U}_r = \\ &\mathbf{U}_r^T E [\mathbf{x}_i \mathbf{x}_i^T] \mathbf{U}_r = \mathbf{U}_r^T \mathbf{U} \Sigma \mathbf{U}^T \mathbf{U}_r = \\ &\Sigma_r = \text{diag}(\sigma_1^2, \dots, \sigma_r^2) \end{aligned}$$

PCA: Minimum MSE reconstruction

- ▶ The projected data is $\mathbf{y}_i = \mathbf{U}_r^T \mathbf{x}_i$
- ▶ If we want to reconstruct the original d -dimensional data from \mathbf{y}_i we calculate

$$\tilde{\mathbf{x}}_i = \mathbf{U}_r \mathbf{y}_i = \mathbf{U}_r \mathbf{U}_r^T \mathbf{x}_i = \mathbf{P}_r \mathbf{x}_i, \quad \text{or} \quad \tilde{\mathbf{X}} = \mathbf{X} \mathbf{U}_r \mathbf{U}_r^T$$

where $\mathbf{P}_r = \mathbf{U}_r \mathbf{U}_r^T$ is a projection matrix in the subspace generated by the r eigenvectors of $\hat{\mathbf{C}}_X$

- ▶ PCA solves the following problem


$$\underset{\mathbf{M}}{\text{minimize}} \sum_{i=1}^n \|\mathbf{x}_i - \tilde{\mathbf{x}}_i\|^2 = \|\mathbf{X} - \mathbf{X} \mathbf{M} \mathbf{M}^T\|_F^2 \quad \text{s.t.} \quad \mathbf{M}^T \mathbf{M} = \mathbf{I}_r$$

whose solution is $\mathbf{M} = \mathbf{U}_r$!!


Example: PCA dimensionality reduction

Iris data set:


- ▶ 3 classes;
- ▶ 150 data;
- ▶ 4 features.



3 principal components:



2 principal components:




Example: Compression/reconstruction with PCA

Reconstructions:

MNIST dataset:

- ▶ 10 classes;
- ▶ 70000 images;
- ▶ 784 pixels.




Example: Compression/reconstruction with PCA


Olivetti Faces dataset:

- ▶ 40 classes;
- ▶ 400 images;
- ▶ 4096 pixels.


Average (“Mean face”):



Principal directions (“eigenfaces”):



Example: Noise reduction with PCA




Using patches of 7x7 pixels.

Source: Shah & Bhalgat (2015). <https://github.com/meetps/CS-663>


Order estimation 1: Scree plot

Plot of the variance explained by each successive principal component:



Order estimation 2: Cumulative explained variance

Cumulative plot of the variance explained by the principal components:



Given a target percentage of explained variance, this plot shows the number of principal components required.

Order estimation techniques

- ▶ Apart from establishing a threshold on the percentage of explained variance, there exist many other methods in the literature to estimate the optimal number of principal components of the model.
- ▶ One of the most popular methods is the **Minimum Description Length** (MDL) criterion [Rissanen, 1978].
- ▶ MDL selects the rank r that minimizes

$$MDL(r) = n \left[\log \left(\prod_{i=1}^r \sigma_i^2 \right) + (d - r) \log \left(\frac{1}{d-r} \sum_{i=r+1}^d \sigma_i^2 \right) \right] + \frac{r(2d-r)+1}{2} \log(n)$$

where n is the total number of observations and d is the dimension of the observation vectors.

Probabilistic PCA (PPCA)

- ▶ So far we have not considered any **probabilistic model** for the input data \mathbf{x}_i
- ▶ Assuming a probabilistic model can help in estimating the most appropriate order or dimension r
- ▶ PPCA (Tipping& Bishop 1999) considers the following generative model of the data

$$\mathbf{x}_i = \mathbf{W}\mathbf{f} + \mathbf{e}$$

with

- ▶ $\mathbf{f} \sim N(0, \mathbf{I}_r)$ a vector of r uncorrelated Gaussian factors (iid)
- ▶ $\mathbf{W} \in \mathcal{R}^{d \times r}$ a deterministic weight matrix
- ▶ $\mathbf{e} \sim N(0, \sigma^2 \mathbf{I}_d)$ a vector of Gaussian iid errors, independent of the factors \mathbf{f}

- ▶ With the above model the observations are also Gaussian $\mathbf{x}_i \sim N(0, \mathbf{C}_x)$ with covariance matrix

$$\mathbf{C}_x = \mathbf{W}\mathbf{W}^T + \sigma^2 \mathbf{I}_d$$

- ▶ From a set of n data we estimate the sample covariance as $\hat{\mathbf{C}}_x = \mathbf{U}\Sigma\mathbf{U}^T$
- ▶ The maximum likelihood (ML) estimates of the parameters of the PPCA model are as follows

$$\begin{aligned}\hat{\mathbf{W}} &= \mathbf{U}_r [\Sigma_r - \hat{\sigma}^2 \mathbf{I}_r]^{1/2} \mathbf{Q} \\ \hat{\sigma}^2 &= \frac{1}{d-r} \sum_{i=r+1}^d \sigma_i^2\end{aligned}$$

where \mathbf{Q} is an arbitrary orthogonal matrix

- ▶ The factors (principal components) can be estimated as follows

$$\hat{\mathbf{f}}_i = \hat{\mathbf{W}}^T \left(\hat{\mathbf{W}}\hat{\mathbf{W}}^T + \hat{\sigma}^2 \mathbf{I}_d \right)^{-1} \mathbf{x}_i$$

Factor Analysis (FA)

- ▶ Spearman 1904
- ▶ Model similar to that assumed in PPCA

$$\mathbf{x}_i = \mathbf{W}\mathbf{f} + \mathbf{e}$$

with the only difference being that the noisess $\mathbf{e} \sim N(0, \mathbf{E})$ can now have different variances in each dimension.

$$\mathbf{E} = \text{diag}(\sigma_1^2, \dots, \sigma_d^2)$$


- ▶ Iterative algorithms are available to obtain ML estimates of the model parameters.

Python exercises

- ▶ **Exercise 2:**
`exercise_2_data_compression_PCA.ipynb`
- ▶ **Exercise 3:**
`exercise_3_digits_mapping_PCA.ipynb`


Linear Discriminant Analysis (LDA)

- ▶ In (supervised) classification problems, the directions of maximum variance (PCA) do not always give features that allow separation between classes.
- ▶ What is the direction of maximum variance in this binary classification problem?




Linear Discriminant Analysis (LDA)


- ▶ In (supervised) classification problems, the directions of maximum variance (PCA) do not always give features that allow separation between classes.
- ▶ What is the direction of maximum variance in this binary classification problem?



► A better direction would be




► A better direction would be



► What is the criterion for finding a projection $y = \mathbf{w}^T \mathbf{x}$ such that the classes are maximally separated?

- A better direction would be



- What is the criterion for finding a projection $y = \mathbf{w}^T \mathbf{x}$ such that the classes are maximally separated?

$$J(\mathbf{w}) = \frac{(m_1 - m_2)^2}{\sigma_1^2 + \sigma_2^2}$$

Problem formulation

LDA or Fisher's Linear Discriminant (Fisher 1936)

LDA obtains the projector \mathbf{w} (of unit norm) that maximizes the function

$$J(\mathbf{w}) = \frac{(m_1 - m_2)^2}{\sigma_1^2 + \sigma_2^2}$$

where (m_i, σ_i^2) are the sample mean and variance of the projected data of the class i

- ▶ Assume a **supervised classification problem** binary with d -dimensional patterns $\mathbf{x}_1, \dots, \mathbf{x}_n$
- ▶ Set of patterns of each class: \mathcal{X}_1 y \mathcal{X}_2
- ▶ $|\mathcal{X}_1| = n_1$ and $|\mathcal{X}_2| = n_2$ with $n_1 + n_2 = n$

- We apply a linear projector \mathbf{w} to the data of the two classes

$$y = \mathbf{w}^T \mathbf{x} \in \mathcal{R}$$

- The set of 1D (projected) patterns of each class we denote as \mathcal{Y}_1 and \mathcal{Y}_2
- The sample mean of the classes in the input and output spaces are

$$\mathbf{m}_i = \frac{1}{n_i} \sum_{i \in \mathcal{X}_i} \mathbf{x}_i, \quad y \quad m_i = \frac{1}{n_i} \sum_{i \in \mathcal{X}_i} \mathbf{w}^T \mathbf{x}_i = \frac{1}{n_i} \sum_{i \in \mathcal{Y}_i} y_i$$

- The distance between the averages is

$$(m_1 - m_2)^2 = \frac{(\mathbf{w}^T \mathbf{m}_1 - \mathbf{w}^T \mathbf{m}_2)^2}{\mathbf{w}^T (\mathbf{m}_1 - \mathbf{m}_2)(\mathbf{m}_1 - \mathbf{m}_2)^T \mathbf{w}}$$

- ▶ Defining the $d \times d$ **between-class scatter matrix**

$$\mathbf{S}_B = (\mathbf{m}_1 - \mathbf{m}_2)(\mathbf{m}_1 - \mathbf{m}_2)^T,$$

the distance is $(m_1 - m_2)^2 = \mathbf{w}^T \mathbf{S}_B \mathbf{w}$

- ▶ The sum of variances in the projected space can be written as

$$\sigma_1^2 + \sigma_2^2 = \frac{1}{n} \mathbf{w}^T \mathbf{S}_W \mathbf{w}$$

where \mathbf{S}_W (**within-class scatter matrix**)

$$\mathbf{S}_W = \left(\sum_{i \in \mathcal{X}_1} (\mathbf{x}_i - \mathbf{m}_1)(\mathbf{x}_i - \mathbf{m}_1)^T + \sum_{i \in \mathcal{X}_2} (\mathbf{x}_i - \mathbf{m}_2)(\mathbf{x}_i - \mathbf{m}_2)^T \right)$$

- Based on the dispersion matrices, the function to be maximized in LDA can be written as follows

$$J(\mathbf{w}) = \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}}$$

which is a generalized Rayleigh quotient.

- Generalized eigenvalue problem

$$\mathbf{S}_B \mathbf{w} = \lambda \mathbf{S}_W \mathbf{w}$$

- The solution for \mathbf{w} is the eigenvector corresponding to the maximum eigenvalue $\mathbf{S}_W^{-1} \mathbf{S}_B$
- In the binary problem under consideration \mathbf{S}_B is of rank one and the solution of the LDA projector is

$$\mathbf{w} = \mathbf{S}_W^{-1} (\mathbf{m}_1 - \mathbf{m}_2)$$

Extensions

- ▶ We can consider more projections by taking more eigenvectors (in this case the projectors are not orthogonal).
- ▶ The extension to the multi-class case with C classes is immediate by defining the inter-class scattering matrix as

$$\mathbf{S}_B = \sum_{j=1}^C n_j (\mathbf{m}_j - \mathbf{m})(\mathbf{m}_j - \mathbf{m})^T$$


where \mathbf{m} is the sample mean of all data, and the intra-class dispersion matrix is

$$\mathbf{S}_W = \sum_{j=1}^C \sum_{\mathbf{x} \in \mathcal{X}_j} (\mathbf{x} - \mathbf{m}_j)(\mathbf{x} - \mathbf{m}_j)^T$$


Example: PCA vs LDA

Iris data set:


- ▶ 3 classes;
- ▶ 150 data;
- ▶ 4 features.



PCA projection,,
2 main components:



LDA projection:




Ejemplo: LDA in supervised classification

Wine dataset:


- ▶ 3 classes;
- ▶ 178 data;
- ▶ 14 features.

Class label	Alcohol	Malic acid	Asa	Alkalinity of ash	Magnesium	Total phenols	Flavonoids	Nonflavonoid phenols	Pteranthocyanidin	Color intensity	OD300/600 of diluted mixtures	Proline		
0	1	14.23	1.71	2.45	38.6	127	2.80	3.06	0.28	2.20	5.04	1.04	3.52	2065
1	1	13.39	1.70	2.14	31.2	300	2.65	2.76	0.28	1.20	4.38	1.05	3.40	1055
2	1	13.18	1.36	2.87	18.6	101	2.80	3.24	0.30	2.81	5.68	1.03	3.17	1183
3	1	14.37	1.85	2.50	38.6	113	3.85	3.49	0.24	2.15	7.00	0.95	3.45	1480
4	1	13.24	2.89	2.87	21.0	116	2.80	2.69	0.38	3.82	4.32	1.04	2.93	739


PCA projection,
2 princ. comp.:




Log. regression
(boundaries):



LDA projection,
2 components:



Log. regression
(boundaries):



Conclusions

- ▶ Dimensionality reduction is a fundamental preprocessing step for many ML techniques.
- ▶ Related to the selection/extraction of features.

Linear dimensionality reduction techniques:

- ▶ **PCA** is the most widely used linear dimensionality reduction technique.
 - ▶ Preprocessing in regression/classification problems; compression/storage of information, etc.
- ▶ Many related techniques. E.g. Linear Discriminant Analysis (**LDA**) which takes into account the class labels.

