

# Dimensionality Reduction: Nonlinear techniques

Minería de Datos (M1966)

Steven Van Vaerenbergh  
Depto. de Matemáticas, Estadística y Computación  
Universidad de Cantabria

Curso 2024-25

Maître Universitario Oficial **Data Science**



# Contents

Multidimensional Scaling

Manifold modeling

Isomap

LLE

t-SNE

Autoencoder

Comparisons

Conclusions

# Nonlinear dimensionality reduction

- ▶ Some data have **structure** that cannot be revealed by linear projections.
- ▶ PCA will fall short.
- ▶ Example: How to make a 2D map of MNIST digits?



- ▶ We need **nonlinear dimensionality reduction** techniques.
- ▶ Concrete benefits include discovering intrinsic structures, and creating a compact visualization that preserves neighborhood relations.

# Multidimensional Scaling (MDS)

- ▶ Torgerson, 1952
- ▶ Nonlinear dimensionality reduction technique originally proposed for visualization in behavioral sciences
- ▶ Data may be of any kind (not just scalar), as long as we can define a **dissimilarity measure** on them:
  - ▶ E.g. respondents are asked to rate similarities between product pairs.
- ▶ MDS seeks a low-dimensional embedding that preserves pairwise dissimilarities as **squared distances**.

# Multidimensional Scaling

1. Define a pairwise dissimilarity measure on the data

$$d_{ij}^* = \text{diss}(\mathbf{x}_i, \mathbf{x}_j)$$

We want to map the data  $\mathbf{x}_i \in \mathbb{R}^d \longrightarrow \mathbf{y}_i \in \mathbb{R}^r$

2. Denote the Euclidean distance on the mapped data as

$$d_{ij} = \|\mathbf{y}_i - \mathbf{y}_j\|$$

3. MDS searches for the mapping that minimizes a cost<sup>1</sup>

$$C = \sum_{i \neq j} (d_{ij}^* - d_{ij})^2$$

---

<sup>1</sup>Classical MDS uses this squared-error (“strain”), while “stress” is another popular variant:  $C' = \sqrt{\sum_{i \neq j} (d_{ij}^* - d_{ij})^2 / \sum_{i \neq j} (d_{ij}^*)^2}$ .

# Classical MDS

Classical MDS assumes  $d^*$  to be Euclidean distances.

1. Matrix of squared dissimilarities  $\mathbf{D}_{ij}^{(2)} = (d_{ij}^*)^2$ .
2. Apply double centering:

$$\mathbf{B} = -\frac{1}{2} \mathbf{J} \mathbf{D}^{(2)} \mathbf{J}$$

$\mathbf{J}$  matrice delle medie

where  $\mathbf{J}$  is the centering matrix  $\mathbf{J} = \mathbf{I} - \mathbf{1}\mathbf{1}^\top/n$ , with  $\mathbf{I}$  the identity matrix and  $\mathbf{1}$  the all-ones vector. Note:  $\mathbf{B} = \mathbf{X}\mathbf{X}^\top$  is the **Gram matrix** (inner products of centered data points).

3. Extract the  $r$  largest **eigenvectors** of  $\mathbf{B}$  and their corresponding eigenvalues; place them in  $\mathbf{E}_r$  and  $\mathbf{\Lambda}_r$ .  
Eigen-decomposition of  $\mathbf{B}$  gives us the directions of maximum variance in this distance-preserving sense.
4. MDS solution:  $\mathbf{Y} = \mathbf{E}_r \mathbf{\Lambda}_r^{1/2}$ .


# MDS example: European cities

We are given a list of distances between cities. We will use MDS to create a map.

City	Barcelona	Belgrade	Berlin	Brussels	Bucharest	...
Barcelona	0	1528.13	1497.61	1062.89	1968.42	...
Belgrade	1528.13	0	999.25	1372.59	447.34	...
Berlin	1497.61	999.25	0	651.62	1293.4	...
Brussels	1062.89	1372.59	651.62	0	1769.69	...
Bucharest	1968.42	447.34	1293.4	1769.69	0	...
⋮	⋮	⋮	⋮	⋮	⋮	⋮

► Python code in `example_3_mds_cities.ipynb`


# MDS example: European cities





# Physical intuition of MDS

1. Randomly place each object on a 2D map;
2. Connect each pair  $(\mathbf{x}_i, \mathbf{x}_j)$  with a spring with the length of the dissimilarity  $d_{ij}^*$ ;
3. Let physics take its course.

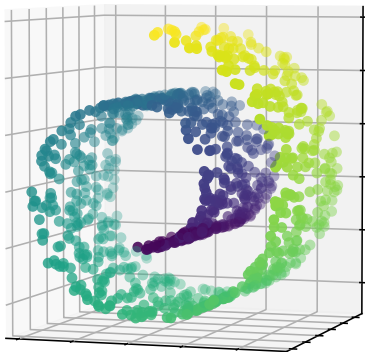


# Landmark MDS


- ▶ MDS requires the eigendecomposition of the dissimilarity matrix, which is computationally expensive.
- ▶ Speedup: “Landmark MDS”
  1. Choose  $q$  data points (“landmarks”), with  $r < q \ll m$ ;
  2. Perform MDS on landmarks;
  3. Map the remaining points to  $\mathbb{R}^r$  using only their distances to landmarks.
- ▶ Landmark MDS combines MDS with the Nyström algorithm for matrix decomposition.
- ▶ By using only  $q$  landmarks, we reduce the memory from  $\mathcal{O}(m^2)$  to  $\mathcal{O}(mq)$ . This is especially important for large datasets where a full matrix decomposition is not feasible.

# Manifold Modeling



- ▶ All previous methods aim to preserve the global distance.
- ▶ What if we want to preserve **local** distance?
- ▶ Example: “Swiss roll” data:



# PCA on Swiss Roll data



# PCA on Swiss Roll data




# Dimensionality Reduction through Manifold Modeling

Manifold: a space that locally resembles Euclidean space.


Main idea: Find a low-dimensional representation that preserves **local properties**.

- ▶ “**local**” → with regard to neighbors.
- ▶ “**properties**” → distances, linear relationships, neighborhood probabilities, etc.





# Manifold modeling

- ▶ In order to preserve the local distance we need to **model the manifold** on which the data lies.
- ▶ Example: “Swiss roll” data:



# Manifold modeling

- ▶ In order to preserve the local distance we need to **model the manifold** on which the data lies.
- ▶ Example: “Swiss roll” data:




- ▶ Local distance on the manifold: **“geodesic distance”**.



# Manifold modeling


- ▶ Swiss roll data: Data on a 2D surface in 3D space.
- ▶ PCA and MDS extract a low-dimensional representation of the data but they do not explicitly model the manifold.
- ▶ PCA and MDS will fail to discover this 2D structure.



- ▶ 2 methods that do model the manifold: Isomap and LLE.


# Isometric feature map (Isomap)

- ▶ Tenenbaum et al., 2000;
- ▶ Idea: Instead of true distance between two points use the **distance along the manifold**;
- ▶ This distance can be very large even if the points are close in  $\mathbb{R}^d$ ;



# Isomap: Steps

1. Construct a **graph** whose nodes are the data points, where a pair of nodes are adjacent only if the two points are close in  $\mathbb{R}^d$ .
2. Approximate the **geodesic distance** along the manifold between any two points as the **shortest path** in the graph;
3. Finally use **MDS** to extract the low-dimensional representation from the resulting matrix of squared distances.



# Dijkstra's algorithm for shortest path calculation

---

**Algorithm 1** Dijkstra's Shortest Path Algorithm.

---

- 1: Input: nodes  $\mathbf{x}_1, \dots, \mathbf{x}_N$  and the distances  $d_{i,j}$  between pairs.
  - 2: **for all**  $i$  **do**
  - 3:     Initialize effective dist.  $\bar{d}_{i,i} = 0$  and  $\bar{d}_{i,j} = \infty$ , for all  $j \neq i$ .
  - 4:     Set node  $i$  as current node  $c$ .
  - 5:     Mark all nodes as "unvisited".
  - 6:     **repeat**
  - 7:         **for all** unvisited nodes  $n$  neighboring current node **do**
  - 8:             **if**  $\bar{d}_{i,n} > \bar{d}_{i,c} + d_{c,n}$  **then**
  - 9:                 Update effective distance:  $\bar{d}_{i,n} = \bar{d}_{i,c} + d_{c,n}$ .
  - 10:             **end if**
  - 11:         **end for**
  - 12:         Mark current node as visited.
  - 13:         Next current = closest unvisited.
  - 14:     **until** all nodes are visited.
  - 15: **end for**
  - 16: Output: effective distances  $\bar{d}_{i,j}$ .
-


# Dijkstra's algorithm for shortest path calculation

---

**Algorithm 2** Dijkstra's Shortest Path Algorithm.


---

- 1: Input: nodes  $\mathbf{x}_1, \dots, \mathbf{x}_N$  and the distances  $d_{i,j}$  between pairs.
- 2: **for all**  $i$  **do**
- 3:   Initialize effective dist.  $\bar{d}_{i,i} = 0$  and  $\bar{d}_{i,j} = \infty$ , for all  $j \neq i$ .
- 4:   Set node  $i$  as current node  $c$ .
- 5:   Mark all nodes as “unvisited”.
- 6:   **repeat**
- 7:     **for all** unvisited nodes  $n$  neighboring current node **do**
- 8:       **if**  $\bar{d}_{i,n} > \bar{d}_{i,c} + d_{c,n}$  **then**
- 9:         Update effective distance:  $\bar{d}_{i,n} = \bar{d}_{i,c} + d_{c,n}$ .
- 10:       **end if**
- 11:     **end for**
- 12:     Mark current node as visited.
- 13:     Next current = closest unvisited.
- 14:   **until** all nodes are visited.
- 15: **end for**
- 16: Output: effective distances  $\bar{d}_{i,j}$ .




# Example: Isomap on Swiss roll data

Data in  $\mathbb{R}^d$



Data in  $\mathbb{R}^2$




- Python code in  
`example_4_isomap_swiss_roll.ipynb`

# Additional properties of Isomap

- ▶ Isomap includes an eigendecomposition, which is **computationally expensive**,  $\mathcal{O}(N)^3$ .
- ▶ Speedup: “Landmark Isomap”, based on Landmark MDS and calculates the geodesic distances only to the landmarks.
- ▶ Isomap does not provide a direct functional form for the mapping  $\mathcal{I} : \mathbb{R}^d \rightarrow \mathbb{R}^r$  that can simply be applied to new data. This further raises computational complexity.

# Locally Linear Embedding (LLE)

- ▶ Roweis & Saul, 2004.
- ▶ Motivation: on a **local** scale the manifold can be approximated by a **linear** subspace.
- ▶ Idea: Model the manifold as a **union of linear patches**.
- ▶ Approximate each point  $\mathbf{x}_i$  as a linear combination of its **neighbors**:  $\mathbf{x}_i \approx \sum_{j \in \mathcal{N}(i)} W_{ij} \mathbf{x}_j$





# LLE: steps

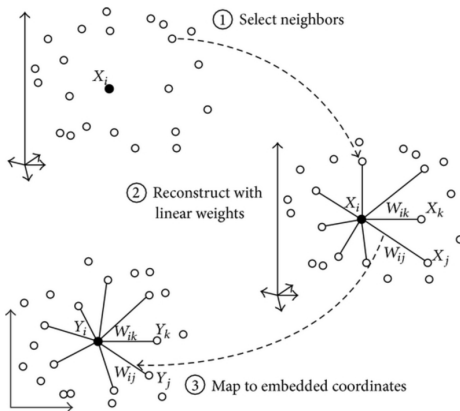
1. Index the nearest neighbors of each  $\mathbf{x}_i \in \mathbb{R}^d$  as  $\mathcal{N}(i)$ .
2. Find the  $W$  that minimizes the reconstruction error

$$\sum_i \|\mathbf{x}_i - \sum_{j \in \mathcal{N}(i)} w_{ij} \mathbf{x}_j\|^2$$

3. Find a set of  $\mathbf{y}_i \in \mathbb{R}^r$  by minimizing

$$\sum_i \|\mathbf{y}_i - \sum_{j \in \mathcal{N}(i)} w_{ij} \mathbf{y}_j\|^2$$

# LLE steps: graphically



# Python exercise

► Exercise 4:

`exercise_4_swiss_roll_reduction_nonlinear.ipynb`

# Extensions: Modified LLE

- ▶ Regularization problem in LLE:
  - ▶ Local  $W$  matrices become **rank-deficient** when the number of **neighbors** is **greater than** the number of input **dimensions**.
  - ▶ Therefore, LLE applies a regularization coefficient.
  - ▶ But regularization distorts the geometry of the manifold.

# Extensions: Modified LLE

- ▶ Regularization problem in LLE:
  - ▶ Local  $W$  matrices become **rank-deficient** when the number of **neighbors** is **greater than** the number of input **dimensions**.
  - ▶ Therefore, LLE applies a regularization coefficient.
  - ▶ But regularization distorts the geometry of the manifold.
- ▶ “Modified” LLE solves this problem by using multiple weight vectors in each neighborhood.
- ▶ Several other extensions: Hessian LLE, LTSA, etc.
- ▶ Note: LLE provides a functional form for the mapping
$$\mathcal{I} : \mathbb{R}^d \rightarrow \mathbb{R}^r$$

# Advanced techniques

- ▶ Popular state-of-the-art techniques:
  - ▶ t-SNE (2008)
  - ▶ UMAP (2018)
- ▶ Capable operating on complex data.
- ▶ Popular in the data science community.

# Stochastic Neighbor Embedding (SNE)

- ▶ Hinton & Roweis, 2003.
- ▶ Constructs a **probability distribution of the potential neighbors** of all  $\mathbf{x}_i$  by placing a Gaussian at each location.
- ▶ Similarly, constructs a probability distribution over all  $\mathbf{y}_i \in \mathbb{R}^r$ .
- ▶ SNE uses gradient descent to **minimize the Kullback-Leibler divergence** between both distributions.
- ▶ Non-convex optimization problem; SNE uses several heuristics.

# t-distributed SNE (t-SNE)

## t-SNE:

- ▶ A recent extension of SNE (van der Maaten & Hinton, 2008).
- ▶ Based on a simpler cost function than SNE and uses Student t-distributions rather than Gaussians.
- ▶ Better results than SNE and faster (converges earlier).
- ▶ Parameter “perplexity”: balance between local and global aspects.
- ▶ Very **flexible** algorithm, but **hard to interpret** and finetune.




# Python exercise


► Exercise 5:

`exercise_5_digits_mapping_nonlinear.ipynb`

# Example: Visualizations of MNIST




Isomap



LLE

# Example: Visualizations of MNIST



# Uniform Manifold Approximation and Projection (UMAP)

- ▶ McInnes & Healy, 2018.
- ▶ Novel manifold learning technique for dimension reduction.
- ▶ Theoretical framework based in Riemannian geometry and algebraic topology.
- ▶ Results comparable to t-SNE, but faster.
- ▶ Python code in `example_5_umap_digits.ipynb`

# Interactive examples


- Isomap: <https://plot.ly/~empet/14345.embed>
- t-SNE: <https://distill.pub/2016/misread-tsne/>

# Interactive examples

- ▶ Isomap: <https://plot.ly/~empet/14345.embed>
- ▶ t-SNE: <https://distill.pub/2016/misread-tsne/>
- ▶ Mapping 6.5M images created by Stable Diffusion  
[https://atlas.nomic.ai/map/  
809ef16a-5b2d-4291-b772-a913f4c8ee61/  
9ed7d171-650b-4526-85bf-3592ee51ea31](https://atlas.nomic.ai/map/809ef16a-5b2d-4291-b772-a913f4c8ee61/9ed7d171-650b-4526-85bf-3592ee51ea31)

# Autoencoder


- ▶ An autoencoder is an **unsupervised** method (typically a neural network) that **compresses** the data to a lower dimension and then **reconstructs** the input back.
- ▶ 2006 seminal paper by Hinton & Salakhutdinov<sup>2</sup>, though idea used since the 1980s.



<sup>2</sup>Hinton, G. E., & Salakhutdinov, R. R. (2006). Reducing the dimensionality of data with neural networks. *Science*, 313(5786), 504-507.

# Autoencoder


- ▶ Does not make specific assumptions about the data.
- ▶ Allows to plug in any regression model as encoder/decoder.
- ▶ If a linear model is used as encoder/decoder, the solution obtained coincides with PCA. (See minimum MSE reconstruction.)



- ▶ Python code in `example_6_autoencoder.ipynb`




# Example: Swiss roll dataset




- S-curve dataset.
- Manifold Learning with 1000 points, 10 neighbors.


MDS (3.5 sec)




Isomap (0.4 sec)




Modified LLE (0.26 sec)



t-SNE (16 sec)




# Example: Spherical dataset




- Spherical data: poles and one slice removed.
- Manifold Learning with 1000 points, 10 neighbors.


MDS (1.7 sec)




Isomap (0.23 sec)




Modified LLE (0.16 sec)



t-SNE (13 sec)




# Example: S-curve dataset




- S-curve dataset.
- Manifold Learning with 1000 points, 10 neighbors.


MDS (3.9 sec)




Isomap (0.47 sec)



Modified LLE (0.32 sec)



t-SNE (21 sec)



# Comparison: Goals

What does each technique try to preserve?

- ▶ PCA: linear structure
- ▶ MDS: global geometry
- ▶ Isomap: geodesic distances (globally)
- ▶ LLE: local translations, rotations, and scalings
- ▶ t-SNE: topology (neighborhood structure)
- ▶ Autoencoder: (generic) encoding that allows for reconstruction

# Comparison: Computational cost

Execution times on 1000 points of 2D toy data:

- ▶ PCA: 0.1 s
- ▶ MDS: 3 s
- ▶ Isomap: 0.2 s (Landmark Isomap)
- ▶ LLE: 0.2 s (Modified LLE)
- ▶ t-SNE: 5 s
- ▶ Autoencoder: 5 s

# Comparison: Restrictions

- ▶ PCA: Unable to discover nonlinear structure.
- ▶ MDS: Requires selection of a meaningful distance.
- ▶ Isomap: Topological instability (affected by noise).
- ▶ LLE: Difficulties to treat manifolds with holes (also the case for Isomap); Requires dense manifold.
- ▶ t-SNE: Slow; hard to interpret; different convergences.
- ▶ Autoencoder: Neural net training requires lots of data.

# Comparison: Guidelines

- ▶ PCA: To decorrelate data; to discover linear structure.
- ▶ MDS: We are given only a distance matrix.
- ▶ Isomap: Data is on a well-connected manifold.
- ▶ LLE: Manifold is approx. linear on a local scale.
- ▶ t-SNE: To visualize complex real-world data.
- ▶ Autoencoder: Dimensionality reduction and reconstruction.

# Other nonlinear dimensionality reduction techniques

- ▶ **Kernel PCA**: nonlinear transformation of data into high-dimensional feature space, then apply PCA in that space.
- ▶ Linear Discriminant Analysis (**LDA**): Supervised dimensionality reduction that maximizes class separation.



# Conclusions

- ▶ Reducing the input data dimensionality is a fundamental preprocessing step for many ML techniques.
- ▶ Related to the selection/extraction of features.
- ▶ Linear dimensionality reduction techniques:
  - ▶ **PCA**: pre-processing for regression/classification techniques; compression/storage of information.
  - ▶ **LDA**: pre-processing in problems of supervised classification.

# Conclusions

- ▶ Nonlinear dimensionality reduction techniques allow to **reveal structure** that linear methods cannot.
- ▶ Higher **computational cost**.
- ▶ Often used for **visualization**, data exploration.
- ▶ Wide range of techniques: **MDS**, **Isomap**, **LLE**, **t-SNE**, **UMAP**, **Autoencoder**, etc.
- ▶ Related: Kernel-based dimensionality reduction techniques (KPCA, KCCA, etc.).

# References

- Christopher J. C. Burges (2010), “Dimension Reduction: A Guided Tour”, Foundations and Trends in Machine Learning: Vol. 2: No. 4, pp 275-365.